

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

a) load/merge data and visualize logerror

```
In [2]: # Load data into DataFrames
df_train = pd.read_csv('train.csv')
df_prop = pd.read_csv('properties.csv')
df_merge = pd.merge(df_train, df_prop, on='id')
```

```
In [3]: # eliminate outliers
percent_range = np.percentile(df_merge.logerror, [1, 99]);
print(df_merge)
df_merge[(df_merge.logerror < percent_range[0])] = df_merge[(df_merge.logerror
df_merge[(df_merge.logerror > percent_range[1])] = df_merge[(df_merge.logerror
```

| | id | logerror | transactiondate | airconditioningtypeid | \ |
|-------|----------|----------|-----------------|-----------------------|---|
| 0 | 14366692 | -0.1684 | 1/1/16 | NaN | |
| 1 | 14739064 | -0.0030 | 1/2/16 | NaN | |
| 2 | 10854446 | 0.3825 | 1/3/16 | NaN | |
| 3 | 11672170 | -0.0161 | 1/3/16 | 1.0 | |
| 4 | 12524288 | -0.0419 | 1/3/16 | NaN | |
| ... | ... | ... | ... | ... | |
| 31720 | 12756771 | 0.0658 | 12/30/16 | NaN | |
| 31721 | 11295458 | -0.0294 | 12/30/16 | 1.0 | |
| 31722 | 11308315 | 0.0070 | 12/30/16 | 1.0 | |
| 31723 | 11703478 | 0.0431 | 12/30/16 | NaN | |
| 31724 | 12566293 | 0.4207 | 12/30/16 | NaN | |

| | architecturalstyletypeid | basementsqft | bathroomcnt | bedroomcnt | \ |
|-------|--------------------------|--------------|-------------|------------|---|
| 0 | NaN | NaN | 3.5 | 4.0 | |
| 1 | NaN | NaN | 1.0 | 2.0 | |
| 2 | NaN | NaN | 2.0 | 2.0 | |
| 3 | NaN | NaN | 4.0 | 5.0 | |
| 4 | NaN | NaN | 1.0 | 1.0 | |
| ... | ... | ... | ... | ... | |
| 31720 | NaN | NaN | 1.0 | 3.0 | |
| 31721 | NaN | NaN | 2.0 | 2.0 | |
| 31722 | NaN | NaN | 3.0 | 5.0 | |
| 31723 | NaN | NaN | 1.0 | 3.0 | |
| 31724 | NaN | NaN | 1.0 | 3.0 | |

| | buildingclasstypeid | buildingqualitytypeid | ... | numberofstories | \ |
|-------|---------------------|-----------------------|-----|-----------------|---|
| 0 | NaN | NaN | ... | NaN | |
| 1 | NaN | NaN | ... | NaN | |
| 2 | NaN | 7.0 | ... | NaN | |
| 3 | NaN | 1.0 | ... | NaN | |
| 4 | NaN | 7.0 | ... | NaN | |
| ... | ... | ... | ... | ... | |
| 31720 | NaN | 7.0 | ... | NaN | |
| 31721 | NaN | 7.0 | ... | NaN | |
| 31722 | NaN | 4.0 | ... | NaN | |
| 31723 | NaN | 7.0 | ... | NaN | |
| 31724 | NaN | 7.0 | ... | NaN | |

| | fireplaceflag | structuretaxvaluedollarcnt | taxvaluedollarcnt | \ |
|-------|---------------|----------------------------|-------------------|---|
| 0 | NaN | 346458.0 | 585529.0 | |
| 1 | NaN | 66834.0 | 210064.0 | |
| 2 | NaN | 55396.0 | 105954.0 | |
| 3 | NaN | 559040.0 | 1090127.0 | |
| 4 | NaN | 56233.0 | 70316.0 | |
| ... | ... | ... | ... | |
| 31720 | NaN | 65728.0 | 307167.0 | |
| 31721 | NaN | 40163.0 | 50203.0 | |
| 31722 | NaN | 248378.0 | 331525.0 | |
| 31723 | NaN | 17520.0 | 39934.0 | |


| | | | | |
|-------|--------------------|-----------------------|-----------|--------------------|
| 31724 | NaN | 66258.0 | 163037.0 | |
| | assessmentyear | landtaxvaluedollarcnt | taxamount | taxdelinquencyflag |
| \ | | | | |
| 0 | 2015 | 239071.0 | 10153.02 | NaN |
| 1 | 2015 | 143230.0 | 2172.88 | NaN |
| 2 | 2015 | 50558.0 | 1443.69 | NaN |
| 3 | 2015 | 531087.0 | 13428.94 | NaN |
| 4 | 2015 | 14083.0 | 913.17 | NaN |
| ... | ... | ... | ... | ... |
| 31720 | 2015 | 241439.0 | 4038.70 | NaN |
| 31721 | 2015 | 10040.0 | 1263.39 | Y |
| 31722 | 2015 | 83147.0 | 6461.79 | NaN |
| 31723 | 2015 | 22414.0 | 627.91 | NaN |
| 31724 | 2015 | 96779.0 | 2560.96 | NaN |
| | taxdelinquencyyear | censustractandblock | | |
| 0 | NaN | NaN | | |
| 1 | NaN | 6.059040e+13 | | |
| 2 | NaN | 6.037140e+13 | | |
| 3 | NaN | 6.037260e+13 | | |
| 4 | NaN | 6.037570e+13 | | |
| ... | ... | ... | | |
| 31720 | NaN | 6.037550e+13 | | |
| 31721 | 15.0 | 6.037900e+13 | | |
| 31722 | NaN | 6.037900e+13 | | |
| 31723 | NaN | 6.037230e+13 | | |
| 31724 | NaN | 6.037540e+13 | | |

[31725 rows x 60 columns]

```
In [4]: # scatter of logerr  
print(df_merge.logerror)  
x = np.arange(0,31725, 1)  
plt.scatter(x, df_merge.logerror,s = 10)  
plt.ylabel('Log Error')  
plt.xlabel('Samples')
```


```
0      -0.1684  
1      -0.0030  
2       0.3825  
3      -0.0161  
4      -0.0419  
...  
31720   0.0658  
31721  -0.0294  
31722   0.0070  
31723   0.0431  
31724   0.4207  
Name: logerror, Length: 31725, dtype: float64
```

Out[4]: Text(0.5, 0, 'Samples')



```
In [5]: # histogram of logerr  
plt.hist(df_merge.logerror)  
plt.xlabel('Log Error')
```

Out[5]: Text(0.5, 0, 'Log Error')



b) data cleaning

```
In [6]: # build new data frame  
missing_vals = df_merge.isna().sum()  
col_num = df_merge.columns.transpose()  
new_df = pd.DataFrame(list(zip(col_num, missing_vals)), columns = ["column_name", "missing_ratio"])  
  
missing_ratio = df_merge.isna().sum() / len(df_merge)  
new_df.insert(2, "missing_ratio", missing_ratio.values)
```

```
In [7]: # fill missing data  
df_merge = df_merge.fillna(df_merge.mean())
```


C:\Users\gandi\AppData\Local\Temp\ipykernel_27800\1960727304.py:2: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
df_merge = df_merge.fillna(df_merge.mean())

c) univariate analysis

```
In [8]: # make bar chart
corr_mat = df_merge.corrwith(df_merge["logerror"])
corr_mat = corr_mat.sort_values()

from matplotlib.colors import TwoSlopeNorm
fig, ax = plt.subplots(figsize=(19, 19))
norm = TwoSlopeNorm(vmin=-1, vcenter = 0, vmax=1)
colors = [plt.cm.RdYlGn(norm(c)) for c in corr_mat.values]
corr_mat.plot.barh(color=colors)
```

Out[8]: <AxesSubplot:>



In [9]: *#Explain*

#The reason behind why some of the values have no correlation values and become NaN is that the standard deviation in the denominator is zero. For these NaN cases, since the denominator is zero, the result of the division is NaN.

d) non-linear regression model

```
In [10]: non_lin_df = df_merge.drop(columns=['id', 'transactiondate', 'hashottuborspa'],
```

```
In [11]: # split and train
from sklearn.model_selection import train_test_split
x_train, x_test = train_test_split(non_lin_df, test_size = 0.3, shuffle = True)

x_train_stats = []
x_test_stats = []
for i in x_train.columns:
    x_train_stats.append([np.mean(x_train[i]), np.std(x_train[i])])
for i in x_test.columns:
    x_test_stats.append([np.mean(x_test[i]), np.std(x_test[i])])

count = 0;
for i in x_train.columns:
    stats = x_train_stats[count]
    if (stats[1] != 0):
        x_train[i] = (x_train[i] - stats[0]) / stats[1]
    count +=1

count = 0;
for i in x_test.columns:
    stats = x_test_stats[count]
    if (stats[1] != 0):
        x_test[i] = (x_test[i] - stats[0]) / stats[1]
    count +=1
from sklearn.neural_network import MLPRegressor as mlp
from sklearn.datasets import make_regression

x_train_new = x_train.drop(["logerror"], axis = 1)
reg = mlp(random_state = 1, max_iter = 500).fit(x_train_new, x_train.logerror)

x_test_new = x_test.drop(['logerror'], axis = 1)
pred = reg.predict(x_test_new)
```

```
In [12]: # report importances and mse
from sklearn.metrics import mean_squared_error
x = mean_squared_error(x_test.logerror, pred)
print("Mean Squared Error: ", x)
```

Mean Squared Error: 1.1200822431414335

Problem 2a.

$$i. \quad \frac{\partial L}{\partial b_k} \quad \hat{y}_k = \frac{e^{b_k}}{\sum_{l=1}^K e^{b_l}} \quad b_k = \sum_{j=0}^2 \beta_{kj} z_j$$

$$L = \sum_{k=1}^K y_k^{(n)} \log(\hat{y}_k^{(n)})$$

$$\frac{\partial L}{\partial b_k} = \sum_{k=1}^K y_k \cdot \frac{\partial \log \hat{y}_k}{\partial b_k} = \sum_{k=1}^K y_k \cdot \frac{1}{\hat{y}_k} \frac{\partial \hat{y}_k}{\partial b_k}$$

$$\hat{y}_k = \frac{e^{b_k}}{\sum_{l=1}^K e^{b_l}} \quad \frac{\partial \hat{y}_k}{\partial b_k} = \hat{y}_k (1 - \hat{y}_k)$$

$$\frac{\partial \hat{y}_k}{\partial b_l} = -\hat{y}_k \hat{y}_l \text{ at } k \neq l$$

$$\begin{aligned} \frac{\partial L}{\partial b_k} &= -y_k^{(n)} (1 - \hat{y}_k) - \sum_{l \neq k} y_l \frac{1}{\hat{y}_l} (-\hat{y}_l \hat{y}_k) \\ &= -y_k^{(n)} (1 - \hat{y}_k) + \sum_{l \neq k} y_l^{(n)} \hat{y}_l \end{aligned}$$

$$\sum_l y_l = 1 \quad \text{not encoded}$$

$$\boxed{\frac{\partial L}{\partial b_k} = \hat{y}_k - y_k}$$

$$ii. \quad \frac{\partial L}{\partial \beta_{kj}} \text{ in terms of } \frac{\partial L}{\partial b_k}$$

$$\frac{\partial L}{\partial \beta_{kj}} = \frac{\partial L}{\partial b_k} \times \frac{\partial b_k}{\partial \beta_{kj}}$$

$$\frac{\partial b_k}{\partial \beta_{kj}} = \frac{\partial}{\partial \beta_{kj}} \sum_{j=0}^2 \beta_{kj} z_j = z_j$$

$$\boxed{\frac{\partial L}{\partial \beta_{kj}} = \frac{\partial L}{\partial b_k} \cdot z_j^T}$$

$$iii. \quad \frac{\partial L}{\partial z} = \frac{\partial L}{\partial b} + B'$$

$$= \frac{\partial L}{\partial b} \cdot \frac{\partial b}{\partial z}$$

$$\boxed{\frac{\partial L}{\partial z} = \frac{\partial L}{\partial b} \cdot B'^T}$$

$$iv. \quad \frac{\partial L}{\partial a_j} \rightarrow \frac{\partial z_i}{\partial a_j} = \frac{e^{-a_j}}{(1+e^{-a_j})^2} \cdot \sum_{i=0}^S x_i^n$$

$$\boxed{\frac{\partial L}{\partial a} = \left(\frac{\partial L}{\partial z} \right)^T X^n \left(\frac{e^{-a}}{(1+e^{-a})^2} \right)^T}$$

Problem 2: Implementing a Multi-layer Perceptron

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
In [2]: #!/pip install seaborn
# Install seaborn (needed to plot confusion matrix) by uncommenting the above
```

```

In [16]: def sigmoid_forward(a):
    # calculates the sigmoid activation function
    # a: pre-activation values
    # returns: activated values

    return 1 / (1 + np.exp(-a));

def sigmoid_backward(grad_accum, a):
    # grad_accum: the gradient of the loss function w.r.t to z
    # a: the pre-activation values
    # returns: the gradient of the loss w.r.t to the preactivation values, a

    return (grad_accum * (np.exp(-a) / ((1 + np.exp(-a)) ** 2)))

def linear_forward(x, weight, bias):
    # Computes the forward pass of the Linear Layer
    # x: input of layer
    # weight, bias: weights and bias of neural network layer
    # returns: output of linear layer

    x = np.column_stack((np.ones(np.shape(x)[0]), x))
    weight = np.column_stack((bias, weight))
    return x @ weight.T

def linear_backward(grad_accum, x, weight, bias):
    # Derivative of the linear layer w.r.t
    # grad_accum: gradient of loss w.r.t function after linear layer
    # returns dl_dw: gradient of loss w.r.t to weights
    # returns dl_dx: gradient of loss w.r.t to input, x
    # return dl_dw, dl_dx

    x = np.column_stack((np.ones(np.shape(x)[0]), x))
    return grad_accum.T @ x, grad_accum @ weight

def softmax_xeloss_forward(b, labels):
    # Input parameters:
    ## b: pre-activation
    # calculates the softmax of the vector b
    # calculates the cross entropy loss between the softmax of b and the label
    # returns: l

    den = np.sum(np.exp(b), axis = 1)
    pred = np.exp(b) / np.tile(den[:, np.newaxis], (1, 10))
    l = -np.sum(labels * (np.log(pred)), axis = 1)
    return l

def softmax_xeloss_backward(yhat, labels):
    # Input parameters:
    # yhat: predictions of the neural network
    # labels: target of the network
    # returns: dl_db gradient of loss w.r.t to b

    dl_db = (-labels * (1 - yhat))
    return dl_db

```

```
def data_load():  
    # load in the data provided in "data/"  
    # Unzip fashion_mnist.zip  
  
    train = np.loadtxt("fashion_mnist/train.csv", delimiter = ",");  
    test = np.loadtxt("fashion_mnist/test.csv", delimiter = ",");  
  
    x_train = train[:, :-1]  
    y_train = train[:, -1]  
  
    x_test = test[:, :-1]  
    y_test = test[:, -1]  
  
    return x_train, y_train, x_test, y_test  
  
def load_params():  
    alpha_weights = np.loadtxt('params/alpha1.txt', delimiter=',')  
    beta_weights = np.loadtxt('params/alpha2.txt', delimiter=',')  
    alpha_bias = np.loadtxt('params/beta1.txt', delimiter=',')  
    beta_bias = np.loadtxt('params/beta2.txt', delimiter=',')  
    return alpha_weights, beta_weights, alpha_bias, beta_bias  
  
def one_hot_encode(y):  
    # convert categorical target features to one hot encoded data  
  
    encode_data = np.zeros((np.shape(y)[0], 10))  
    y = np.array(y, dtype = "int")  
    for column in range(np.shape(y)[0]):  
        encode_data[column, y[column]] = 1  
    return encode_data  
  
def train(batchsize=1 , eta = 0.01, num_epochs=100, h = 256, init='default'):  
    x_train, y_train, x_test, y_test = data_load()  
  
    y_train = one_hot_encode(y_train)  
    y_test = one_hot_encode(y_test)  
  
    if init == 'default':  
        alpha_weights, beta_weights, alpha_bias, beta_bias = load_params()  
    elif init=='zeros':  
        # initialize weights and biases to 0  
        alpha_weights, beta_weights, alpha_bias, beta_bias = load_params() * 0  
    elif init=='ones':  
        # initialize weights and biases to 1  
        alpha_weights, beta_weights, alpha_bias, beta_bias = load_params() * 0  
    elif init=='random':  
        # initialize weights and biases to random values between -1 and 1  
        pass  
  
    train_loss_list = []  
    test_loss_list = []  
    acc_list = []  
  
    for epoch in (range(num_epochs)):  
        print("Epoch :", epoch)
```

```

for batch in range(int(len(x_train) / batchsize) + (len(x_train) % batchsize) != 0):
    batch_x = x_train[batch * batchsize:(batch + 1) * batchsize, :]
    batch_y = y_train[batch * batchsize:(batch + 1) * batchsize, :]

    ##### FORWARD
    # Linear -> Sigmoid -> Linear -> Softmax
    A = linear_forward(batch_x, alpha_weights, alpha_bias)
    Z = sigmoid_forward(A)
    B = linear_forward(Z, beta_weights, beta_bias)
    L = softmax_xeloss_forward(B, batch_y)

    ##### BACKWARD
    den = np.sum(np.exp(B), axis = 1)
    yhat = np.exp(B) / den
    dl_db = softmax_xeloss_backward(yhat, batch_y)
    dl_dbeta, dl_dz = linear_backward(dl_db, Z, beta_weights, beta_bias)
    dl_da = sigmoid_backward(dl_dz, A)
    dl_dalpha, dl_dx = linear_backward(dl_da, batch_x, alpha_weights, alpha_bias)

    ##### UPDATE
    alpha_weights = alpha_weights - dl_dalpha[:, 1:] * eta
    beta_weights = beta_weights - dl_dbeta[:, 1:] * eta
    alpha_bias = alpha_bias - dl_dalpha[:, 0] * eta
    beta_bias = beta_bias - dl_dbeta[:, 0] * eta

    # store average training loss for the epoch
    # calculate test predictions and loss
    num_test = np.shape(x_train)[0]
    A = linear_forward(x_train, alpha_weights, alpha_bias)
    Z = sigmoid_forward(A)
    B = linear_forward(Z, beta_weights, beta_bias)
    L = softmax_xeloss_forward(B, y_train)

    train_loss_list.append(np.sum(L)/num_test)
    y_hat_train = np.exp(B)/np.sum(np.exp(B), axis = 0)
    train_pred = np.argmax(y_hat_train, axis = 1)[:]
    vector_corr = train_pred == np.argmax(y_train, axis = 1)[:]

    #Test
    num_test = np.shape(x_test)[0]
    A = linear_forward(x_test, alpha_weights, alpha_bias)
    Z = sigmoid_forward(A)
    B = linear_forward(Z, beta_weights, beta_bias)
    L = softmax_xeloss_forward(B, y_test)

    test_loss_list.append(np.sum(L)/num_test)
    den = np.sum(np.exp(B), axis = 1)
    y_hat_test = np.exp(B)/np.sum(np.exp(B), axis = 0)
    test_pred = np.exp(B)/np.sum(np.exp(B), axis = 0)
    vector_correct = np.argmax(test_pred, axis = 1)[:] == np.argmax(y_test, axis = 1)[:]

    # calculate test accuracy
    acc_list.append(np.sum(vector_correct) / num_test)

# return train_loss_list, test_loss_list, as well as test and train predictions
return train_loss_list, test_loss_list, acc_list, y_hat_train, y_hat_test

```

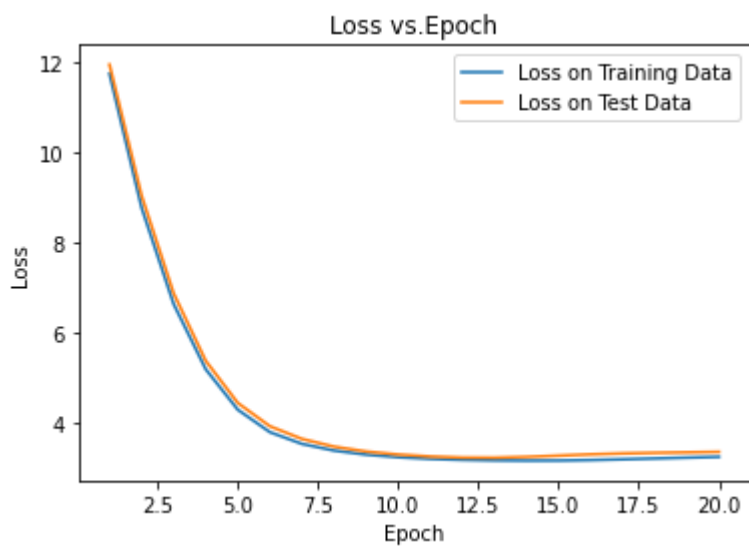
```
In [17]: train_loss_list, test_loss_list, acc_list, yhatTrain, yhatTest = train(num_epo
```

```
Epoch : 0  
Epoch : 1  
Epoch : 2  
Epoch : 3  
Epoch : 4  
Epoch : 5  
Epoch : 6  
Epoch : 7  
Epoch : 8  
Epoch : 9  
Epoch : 10  
Epoch : 11  
Epoch : 12  
Epoch : 13  
Epoch : 14  
Epoch : 15  
Epoch : 16  
Epoch : 17  
Epoch : 18  
Epoch : 19
```

Plot Loss

```
In [18]: # Plot training loss, testing loss as a function of epochs
```

```
In [19]: epochs = np.arange(1, 21);  
plt.figure()  
plt.plot(epochs, train_loss_list, label = "Loss on Training Data")  
plt.plot(epochs, test_loss_list, label = "Loss on Test Data")  
plt.title("Loss vs.Epoch")  
plt.ylabel("Loss")  
plt.xlabel("Epoch")  
plt.legend()  
plt.show()
```



Confusion Matrix

```
In [20]: def plot_confusion(yhat, y, title = '[Training or Test] Set'):

    pred_train = np.argmax(yhat, axis=1)
    true_train = np.argmax(y, axis=1)
    print(true_train.shape)
    conf_train = np.zeros((10,10))
    for i in range(len(y)):
        conf_train[ true_train[i], pred_train[i] ] += int(1)

    sns.heatmap(conf_train, annot=True, fmt='.3g')
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Title')
    plt.show()

# plot_confusion(yhat_train, y_train, title = "Training Set")
# plot_confusion(yhat_test, y_test, title = "Test Set")
#yhat: predictions
#y: one-hot-encoded labels

X_train, y_train, X_test, y_test = data_load()
y_test = one_hot_encode(y_test)
plot_confusion(yhatTest, y_test, title = "Test Data")

(1000,)
```



Correct and Incorrect Classification Samples

```
In [8]: def plot_image(vector, out_f_name, label=None):  
        """  
        Takes a vector as input of size (784) and saves as an image  
        """  
        image = np.asarray(vector).reshape(28, 28)  
        plt.imshow(image, cmap='gray')  
        if label:  
            plt.title(label)  
        plt.axis('off')  
        plt.savefig(f'{out_f_name}.png', bbox_inches='tight')  
        plt.show()
```

```
In [9]: # Use plot_image function to display samples that are correctly and incorrectl
```

Effect Of Learning Rate

```
In [10]: # Plot test loss as a function of epochs
```

Effect of Initialization

```
In [11]: # Plot test loss as a function of epochs
```


Question 3: CIFAR-10 Classification using CNN

- Please **do not** change the default variable names in this problem, as we will use them in different parts.
- The default variables are initially set to "None".

```
In [1]: import numpy as np # linear algebra
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torchvision
from torchvision import datasets, transforms, models
from torch.utils.data import *
import random
from tqdm import tqdm
import warnings
```

```
In [2]: def imshow(img):
    img = img / 2 + 0.5 # unnormalize
    npimg = img.numpy() # convert from tensor
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()
```

In [3]:

```
# The below two lines are optional and are just there to avoid any SSL
# related errors while downloading the CIFAR-10 dataset
import ssl
ssl._create_default_https_context = ssl._create_unverified_context

#Initializing normalizing transform for the dataset
normalize_transform = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize(mean = (0.5, 0.5, 0.5),
                                     std = (0.5, 0.5, 0.5))])

#Downloading the CIFAR10 dataset into train and test sets
train_dataset = torchvision.datasets.CIFAR10(
    root="./CIFAR10/train", train=True,
    transform=normalize_transform,
    download=True)

test_dataset = torchvision.datasets.CIFAR10(
    root="./CIFAR10/test", train=False,
    transform=normalize_transform,
    download=True)

#Generating data loaders from the corresponding datasets
batch_size = 128
train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size)

classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog',
           'frog', 'horse', 'ship', 'truck')

# get first 100 training images
dataiter = iter(train_loader)
imgs, lbls = dataiter.next()

for i in range(20):
    plt.title(classes[lbls[i]])
    imshow(imgs[i])
```

Files already downloaded and verified
Files already downloaded and verified



```
In [4]: # check pytorch cuda and use cuda if possible
device = torch.cuda.is_available()
print('*' * 50)
if torch.cuda.is_available():
    print('CUDA is found! Training on %s.....'%torch.cuda.get_device_name(0))
else:
    warnings.warn('CUDA not found! Training may be slow.....')
```

CUDA is found! Training on NVIDIA GeForce RTX 2070 with Max-Q Design.....

P1. Build you own CNN model

TODO

- Design your model class in **CNNModel(nn.Module)** and write forward pass in **forward(self, x)**
- Create loss function in **error**, optimizer in **optimizer**
- Define hyperparameters: **learning_rate**, **num_epochs**
- Plot your **loss vs num_epochs** and **accuracy vs num_epochs**

Hints

- Start with low number of epochs for debugging. (eg. num_epochs=1)
- Be careful with the input dimension of fully connected layer.
- The dimension calculation of the output tensor from the input tensor is \
$$D_{out} = \frac{D_{in} - K + 2P}{S} + 1$$

$$D_{out} : \text{Dimension of output tensor} \quad D_{in} : \text{Dimension of input tensor} \quad K : \text{width/height of the kernel} \quad S : \text{stride} \quad P : \text{padding}$$

Convolutional and Pooling Layers

A convolutional layer using pyTorch:

```
torch.nn.Conv2d(num_in_channels, num_out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

For example:

```
torch.nn.Conv2d(3, 32, 3)
```

It applies a 2D convolution over an input signal composed of several input planes. If we have

input size with (N, C_{in}, H, W) and output size with $(N, C_{out}, H_{out}, W_{out})$, the 2D convolution can be described as

$$out(N_i, C_{out_j}) = bias(C_{out_j}) + \sum_{k=0}^{C_{in}-1} weight(C_{out_j}, k) \star input(N_i, k)$$

num_in_channels: is the number of channels of the input tensor. If the previous layer is the input layer, num_in_channels is the number of channels of the image (3 channels for RGB images), otherwise num_in_channels is equal to the number of feature maps of the previous layer.

num_out_channels: is the number of filters (feature extractor) that this layer will apply over the image or feature maps generated by the previous layer.

kernel_size: is the size of the convolving kernel

stride: is the stride of the convolution. Default: 1

padding: is the padding added to all four sides of the input. Default: 0

dilation: is the spacing between kernel elements. Default: 1

group: is the number of blocked connections from input channels to output channels. Default: 1

bias: If True, adds a learnable bias to the output. Default: True

A Simple Convolutional Neural Network

In our convnet we'll initially use this structure shown below:

*input -> convolution -> fully connected -> output *

At the end of the last convolutional layer, we get a tensor of dimension (num_channels, height, width). Since now we are going to feed it to a fully connected layer, we need to convert it into a 1-D vector, and for that we use the reshape method:

```
x = x.view(x.size(0), -1)
```

The way of calculating size of the output size from previous convolution layer can be formulized as below:

$$H_{output} = \frac{H_{in} + 2 \times padding - kernel_Size}{stride} + 1$$

For more details, you can refer to this link: \ <https://pytorch.org/docs/stable/generated>

```
In [5]: class CNNModel(nn.Module):
def __init__(self):
    super(CNNModel, self).__init__()
    # TODO: Create CNNModel using 2D convolution. You should vary the number of
    # In this function, you should define each of the individual components of
    # Example:
    # self.cnn1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=5, str
    # self.relu1 = nn.ReLU()
    # self.maxpool1 = nn.MaxPool2d(kernel_size=2)
    self.cnn1 = torch.nn.Conv2d(in_channels = 3, out_channels = 6, kernel_size
    self.relu = torch.nn.ReLU()
    self.cnn2 = torch.nn.Conv2d(in_channels = 6, out_channels = 16, kernel_siz
    self.cnn3 = torch.nn.Conv2d(in_channels = 16, out_channels = 24, kernel_si

    # TODO: Create Fully connected layers. You should calculate the dimension
    # Example:
    # self.fc1 = nn.Linear(16 * 110 * 110, 5)
    # Fully connected 1
    self.fc1 = torch.nn.Linear(24*24*24, 120)
    self.fc2 = torch.nn.Linear(120, 84)
    self.fc3 = torch.nn.Linear(84, 10)

def forward(self,x):

    # TODO: Perform forward pass in below section
    # In this function, you will apply the components defined earlier to the i
    # Example:
    out = self.cnn1(x)
    out = self.relu(out)
    out = self.cnn2(out)
    out = self.relu(out)
    #plt.imshow(out[0][0].cpu().detach().numpy())
    #plt.show()
    #plt.close('all')
    out = self.cnn3(out)
    out = self.relu(out)
    #out = self.relu1(out)
    # out = self.maxpool1(out)
    # to visualize feature map in part a, part b.i), use the following three l

    out = out.view(out.size(0), -1)
    out = self.fc1(out)
    out = self.relu(out)
    out = self.fc2(out)
    out = self.relu(out)
    out = self.fc3(out)
    return out
```

Starting Up Our Model

We'll send the model to our GPU if you have one so we need to create a CUDA device and instantiate our model. Then we will define the loss function and hyperparameters that we need to train the model: \

###TODO

- Define Cross Entropy Loss
- Create Adam Optimizer
- Define hyperparameters

```
In [6]: # Create CNN
device = "cuda" if torch.cuda.is_available() else "cpu"
model = CNNModel()
model.to(device)

# TODO: define Cross Entropy Loss
error = torch.nn.CrossEntropyLoss()

# TODO: create Adam Optimizer and define your hyperparameters
learning_rate = 0.001
optimizer = torch.optim.Adam(model.parameters(), learning_rate)
num_epochs = 20

from torchsummary import summary
batch_size = 16
summary(model, input_size=(3, 32, 32))
```

```
=====
Layer (type:depth-idx)                   Param #
=====
|---Conv2d: 1-1                          456
|---ReLU: 1-2                            --
|---Conv2d: 1-3                          880
|---Conv2d: 1-4                          3,480
|---Linear: 1-5                           1,659,000
|---Linear: 1-6                           10,164
|---Linear: 1-7                           850
=====
Total params: 1,674,830
Trainable params: 1,674,830
Non-trainable params: 0
=====
```

```
Out[6]: =====
Layer (type:depth-idx)                   Param #
=====
|---Conv2d: 1-1                          456
|---ReLU: 1-2                            --
|---Conv2d: 1-3                          880
|---Conv2d: 1-4                          3,480
|---Linear: 1-5                           1,659,000
|---Linear: 1-6                           10,164
|---Linear: 1-7                           850
=====
Total params: 1,674,830
Trainable params: 1,674,830
Non-trainable params: 0
=====
```

Training the Model

TODO

- Make predictions from your model
- Calculate Cross Entropy Loss from predictions and labels

```
In [7]: count = 0
loss_list = []
iteration_list = []
accuracy_list = []
for epoch in tqdm(range(num_epochs)):
    model.train()
    for i, (images, labels) in enumerate(train_loader):
        images, labels = images.to(device), labels.to(device)

        # Clear gradients
        optimizer.zero_grad()

        # TODO: Forward propagation
        outputs = model(images)

        # TODO: Calculate softmax and cross entropy loss
        loss = error(outputs, labels)

        # Backpropagate your Loss
        torch.sum(loss).backward()

        # Update CNN model
        optimizer.step()

    count += 1

    if count % 50 == 0:
        model.eval()
        # Calculate Accuracy
        correct = 0
        total = 0
        # Iterate through test dataset
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)

            # Forward propagation
            outputs = model(images)

            # Get predictions from the maximum value
            predicted = torch.argmax(outputs,1)

            # Total number of labels
            total += len(labels)

            correct += (predicted == labels).sum()

        accuracy = 100 * correct / float(total)

        # store loss and iteration
        loss_list.append(loss.item())
        iteration_list.append(count)
        accuracy_list.append(accuracy.item())
    if count % 500 == 0:
        # Print Loss
        print('Iteration: {} Loss: {} Accuracy: {} %'.format(count, loss
```



```
5%|██████████| 1/20 [00:39<12:22, 39.06s/it]
Iteration: 500 Loss: 1.3727850914001465 Accuracy: 51.8599967956543 %

10%|██████████| 2/20 [01:14<11:01, 36.77s/it]
Iteration: 1000 Loss: 1.1488912105560303 Accuracy: 59.34000015258789 %

15%|██████████| 3/20 [01:51<10:32, 37.19s/it]
Iteration: 1500 Loss: 0.8596102595329285 Accuracy: 61.46999740600586 %

25%|██████████| 5/20 [03:04<09:10, 36.67s/it]
Iteration: 2000 Loss: 0.941204845905304 Accuracy: 61.87999725341797 %

30%|██████████| 6/20 [03:37<08:17, 35.53s/it]
Iteration: 2500 Loss: 0.83780837059021 Accuracy: 62.68000030517578 %

35%|██████████| 7/20 [04:12<07:37, 35.20s/it]
Iteration: 3000 Loss: 0.8059139847755432 Accuracy: 61.82999801635742 %

40%|██████████| 8/20 [04:49<07:08, 35.74s/it]
Iteration: 3500 Loss: 0.5071962475776672 Accuracy: 57.68000030517578 %

50%|██████████| 10/20 [06:02<06:01, 36.18s/it]
Iteration: 4000 Loss: 0.3492661714553833 Accuracy: 60.43000030517578 %

55%|██████████| 11/20 [06:40<05:31, 36.88s/it]
Iteration: 4500 Loss: 0.37602102756500244 Accuracy: 58.279998779296875 %

60%|██████████| 12/20 [07:14<04:47, 35.97s/it]
Iteration: 5000 Loss: 0.35433849692344666 Accuracy: 56.59000015258789 %

70%|██████████| 14/20 [08:26<03:35, 35.91s/it]
Iteration: 5500 Loss: 0.19080926477909088 Accuracy: 58.25 %

75%|██████████| 15/20 [09:02<02:59, 36.00s/it]
Iteration: 6000 Loss: 0.1000833660364151 Accuracy: 58.459999084472656 %

80%|██████████| 16/20 [09:38<02:23, 35.87s/it]
Iteration: 6500 Loss: 0.2079406976699829 Accuracy: 57.599998474121094 %

85%|██████████| 17/20 [10:09<01:43, 34.49s/it]
Iteration: 7000 Loss: 0.11518188565969467 Accuracy: 58.519996643066406 %

95%|██████████| 19/20 [11:28<00:36, 36.95s/it]
Iteration: 7500 Loss: 0.23340679705142975 Accuracy: 55.98999786376953 %



100%|██████████| 20/20 [12:04<00:00, 36.23s/it]
```

```
In [8]: # !pip install torchsummary
# uncomment above line if you don't have torchsummary installed yet
# Print torchsummary of model
from torchsummary import summary
print(summary(model, input_size = (3, 32, 32)))
print("Kernel Size: 3x3")
```

```
=====
Layer (type:depth-idx)                   Param #
=====
├─Conv2d: 1-1                             456
├─ReLU: 1-2                               --
├─Conv2d: 1-3                             880
├─Conv2d: 1-4                             3,480
├─Linear: 1-5                             1,659,000
├─Linear: 1-6                             10,164
├─Linear: 1-7                             850
=====
Total params: 1,674,830
Trainable params: 1,674,830
Non-trainable params: 0
=====
Layer (type:depth-idx)                   Param #
=====
├─Conv2d: 1-1                             456
├─ReLU: 1-2                               --
├─Conv2d: 1-3                             880
├─Conv2d: 1-4                             3,480
├─Linear: 1-5                             1,659,000
├─Linear: 1-6                             10,164
├─Linear: 1-7                             850
=====
Total params: 1,674,830
Trainable params: 1,674,830
Non-trainable params: 0
=====
Kernel Size: 3x3
```

```
In [9]: # visualization loss
plt.plot(iteration_list,loss_list)
plt.xlabel("Number of iteration")
plt.ylabel("Loss")
plt.title("CNN: Loss vs Number of iteration")
plt.show()

# visualization accuracy
plt.plot(iteration_list,accuracy_list,color = "red")
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("CNN: Accuracy vs Number of iteration")
plt.show()
```



Evaluating the Model


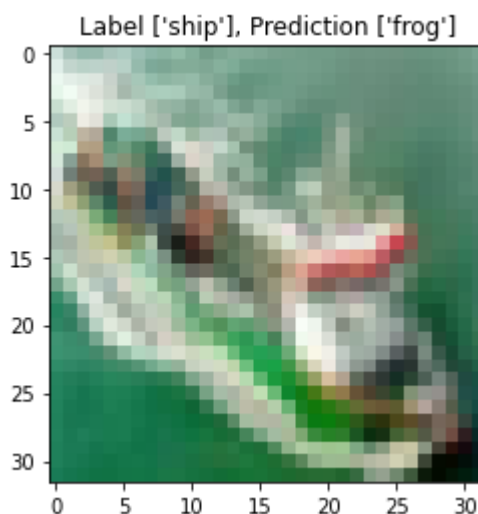
```
In [10]: import random
#To-do: evaluate on test set, instead of training set
random_image = random.randint(0,len(test_dataset))
image = test_dataset.__getitem__(random_image)
model.eval()
images, labels = next(iter(test_loader))
images, labels = images.to(device), labels.to(device)
predictions = torch.argmax(model(images),1)
num_cols=1
num_rows = 25# Len(labels)
label_map = [['airplane'],['automobile'],['bird'],['cat'], ['deer'], ['dog'],

r = list(range(num_rows))
random.shuffle(r)
rand_range = r[0:5]
for idx in rand_range:
    img = images.cpu()[idx]

    plt.title(f"Label {label_map[labels[idx]]}, Prediction {label_map[predictions[idx]]}")
    imshow(img)

    plt.axis('off')

plt.show()
```



Label ['airplane'], Prediction ['airplane']
Label ['dog'], Prediction ['deer']



Label ['horse'], Prediction ['horse']




```
In [11]: #To-do: evaluate on test set, instead of training set
random_image = random.randint(0, len(test_dataset))
image = test_dataset.__getitem__(random_image)
model.eval()
images, labels = next(iter(test_loader))
images, labels = images.to(device), labels.to(device)
predictions = torch.argmax(model(images), 1)
num_cols=1
num_rows = 25# len(labels)
label_map = [['airplane'], ['automobile'], ['bird'], ['cat'], ['deer'], ['dog'],

for idx in range(num_rows):
    img = images.cpu()[idx]

    plt.title(f"Label {label_map[labels[idx]]}, Prediction {label_map[predictions[idx]]}")
    imshow(img)

    plt.axis('off')
plt.show()
```



Part 2: Additional Components

i.)

```

In [12]: class CNNModel(nn.Module):
def __init__(self):
    super(CNNModel, self).__init__()
    # TODO: Create CNNModel using 2D convolution. You should vary the number of
    # In this function, you should define each of the individual components of
    # Example:
    # self.cnn1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=5, str
    # self.relu1 = nn.ReLU()
    # self.maxpool1 = nn.MaxPool2d(kernel_size=2)
    self.cnn1 = torch.nn.Conv2d(in_channels = 3, out_channels = 6, kernel_size
    self.relu = torch.nn.ReLU()
    self.cnn2 = torch.nn.Conv2d(in_channels = 6, out_channels = 16, kernel_siz

    self.maxpool1 = nn.MaxPool2d(2)

    self.cnn3 = torch.nn.Conv2d(in_channels = 16, out_channels = 24, kernel_si
    self.cnn4 = torch.nn.Conv2d(in_channels = 24, out_channels = 24, kernel_si
    self.cnn5 = torch.nn.Conv2d(in_channels = 24, out_channels = 24, kernel_si

    # TODO: Create Fully connected layers. You should calculate the dimension
    # Example:
    # self.fc1 = nn.Linear(16 * 110 * 110, 5)
    # Fully connected 1
    self.fc1 = torch.nn.Linear(13824, 120)
    self.fc2 = torch.nn.Linear(120, 84)
    self.fc3 = torch.nn.Linear(84, 10)

def forward(self,x):

    # TODO: Perform forward pass in below section
    # In this function, you will apply the components defined earlier to the i
    # Example:
    out = self.cnn1(x)
    out = self.relu(out)
    out = self.cnn2(out)
    out = self.relu(out)
    #plt.imshow(out[0][0].cpu().detach().numpy())
    #plt.show()
    #plt.close('all')
    out = self.cnn3(out)
    out = self.relu(out)
    #out = self.relu1(out)
    # out = self.maxpool1(out)
    # to visualize feature map in part a, part b.i), use the following three l

    out = out.view(out.size(0), -1)
    out = self.fc1(out)
    out = self.relu(out)
    out = self.fc2(out)
    out = self.relu(out)
    out = self.fc3(out)
    return out

# Create CNN
device = "cuda" if torch.cuda.is_available() else "cpu"

```

```
model = CNNModel()
model.to(device)

# TODO: define Cross Entropy Loss
error = torch.nn.CrossEntropyLoss()

# TODO: create Adam Optimizer and define your hyperparameters
learning_rate = 1e-3
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
num_epochs = 20

from torchsummary import summary
print(summary(model, input_size=(3, 32, 32)))

count = 0
loss_list_bi = []
iteration_list_bi = []
accuracy_list_bi = []
for epoch in tqdm(range(num_epochs)):
    model.train()
    for i, (images, labels) in enumerate(train_loader):
        images, labels = images.to(device), labels.to(device)

        # Clear gradients
        optimizer.zero_grad()

        # TODO: Forward propagation
        outputs = model(images)

        # TODO: Calculate softmax and cross entropy Loss
        loss = error(outputs, labels)

        # Backpropagate your Loss
        T = torch.sum(loss)
        T.backward()

        # Update CNN model
        optimizer.step()

    count += 1

    if count % 50 == 0:
        model.eval()
        # Calculate Accuracy
        correct = 0
        total = 0
        # Iterate through test dataset
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)

            # Forward propagation
            outputs = model(images)

            # Get predictions from the maximum value
            predicted = torch.argmax(outputs, 1)
```



```

        # Total number of labels
        total = total + len(labels)
        correct = correct + (predicted == labels).sum()

    accuracy = 100 * correct / float(total)

    # store loss and iteration
    loss_list_bi.append(loss.item())
    iteration_list_bi.append(count)
    accuracy_list_bi.append(accuracy.item())
if count % 500 == 0:
    # Print Loss
    print('Iteration: {} Loss: {} Accuracy: {} %'.format(count, loss

```

```

=====
Layer (type:depth-idx)                   Param #
=====
├─Conv2d: 1-1                             456
├─ReLU: 1-2                               --
├─Conv2d: 1-3                             880
├─MaxPool2d: 1-4                          --
├─Conv2d: 1-5                             3,480
├─Conv2d: 1-6                             5,208
├─Conv2d: 1-7                             5,208
├─Linear: 1-8                             1,659,000
├─Linear: 1-9                             10,164
├─Linear: 1-10                            850
=====
Total params: 1,685,246
Trainable params: 1,685,246
Non-trainable params: 0
=====
Layer (type:depth-idx)                   Param #
=====
├─Conv2d: 1-1                             456
├─ReLU: 1-2                               --
├─Conv2d: 1-3                             880
├─MaxPool2d: 1-4                          --
├─Conv2d: 1-5                             3,480
├─Conv2d: 1-6                             5,208
├─Conv2d: 1-7                             5,208
├─Linear: 1-8                             1,659,000
├─Linear: 1-9                             10,164
├─Linear: 1-10                            850
=====
Total params: 1,685,246
Trainable params: 1,685,246
Non-trainable params: 0
=====

```

```
5%|█          | 1/20 [00:31<10:03, 31.79s/it]
```

```
Iteration: 500 Loss: 1.284759521484375 Accuracy: 50.55999755859375 %
```



```
10%|█         | 2/20 [01:08<10:23, 34.62s/it]
```

```
Iteration: 1000 Loss: 1.1590650081634521 Accuracy: 57.529998779296875 %  
15%|███████| 3/20 [01:43<09:51, 34.81s/it]  
Iteration: 1500 Loss: 0.8989694714546204 Accuracy: 60.04999923706055 %  
25%|███████| 5/20 [02:55<08:55, 35.72s/it]  
Iteration: 2000 Loss: 0.8936105370521545 Accuracy: 61.47999954223633 %  
30%|███████| 6/20 [03:27<08:03, 34.54s/it]  
Iteration: 2500 Loss: 0.6931913495063782 Accuracy: 61.0099983215332 %  
35%|███████| 7/20 [04:03<07:31, 34.74s/it]  
Iteration: 3000 Loss: 0.816838264465332 Accuracy: 58.89999771118164 %  
40%|███████| 8/20 [04:37<06:54, 34.55s/it]  
Iteration: 3500 Loss: 0.6166757941246033 Accuracy: 58.31999969482422 %  
50%|███████| 10/20 [05:43<05:38, 33.87s/it]  
Iteration: 4000 Loss: 0.6007229089736938 Accuracy: 58.769996643066406 %  
55%|███████| 11/20 [06:16<05:00, 33.39s/it]  
Iteration: 4500 Loss: 0.5166906118392944 Accuracy: 57.41999816894531 %  
60%|███████| 12/20 [06:46<04:19, 32.43s/it]  
Iteration: 5000 Loss: 0.31766220927238464 Accuracy: 57.119998931884766 %  
70%|███████| 14/20 [07:54<03:20, 33.37s/it]  
Iteration: 5500 Loss: 0.25976741313934326 Accuracy: 56.5 %  
75%|███████| 15/20 [08:28<02:46, 33.32s/it]  
Iteration: 6000 Loss: 0.3657153844833374 Accuracy: 56.05999755859375 %  
80%|███████| 16/20 [09:01<02:13, 33.48s/it]  
Iteration: 6500 Loss: 0.27545157074928284 Accuracy: 57.73999786376953 %  
85%|███████| 17/20 [09:33<01:38, 32.99s/it]  
Iteration: 7000 Loss: 0.1770372986793518 Accuracy: 56.87999725341797 %  
95%|███████| 19/20 [10:41<00:33, 33.41s/it]  
Iteration: 7500 Loss: 0.1567377895116806 Accuracy: 57.439998626708984 %  
100%|███████| 20/20 [11:14<00:00, 33.73s/it]
```

```
In [13]: # visualization loss
plt.plot(iteration_list_bi,loss_list_bi, color = "blue", label = 'Maxpool')
plt.plot(iteration_list,loss_list, color = "red", label = 'Original')
plt.xlabel("Number of iteration")
plt.ylabel("Loss")
plt.title("CNN: Loss vs Number of iteration")
plt.legend()
plt.show()

# visualization accuracy
plt.plot(iteration_list,accuracy_list,color = "red", label = 'Original')
plt.plot(iteration_list_bi,accuracy_list_bi,color = "blue", label = 'Maxpool')
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("CNN: Accuracy vs Number of iteration")
plt.legend()
plt.show()

print("\nMemory Required for Model: 6.85 MB")
print("\nIn terms of observations, it seems that the Maxpool and Original meth
```



Memory Required for Model: 6.85 MB

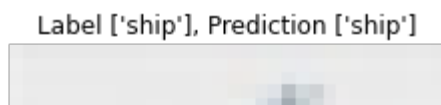
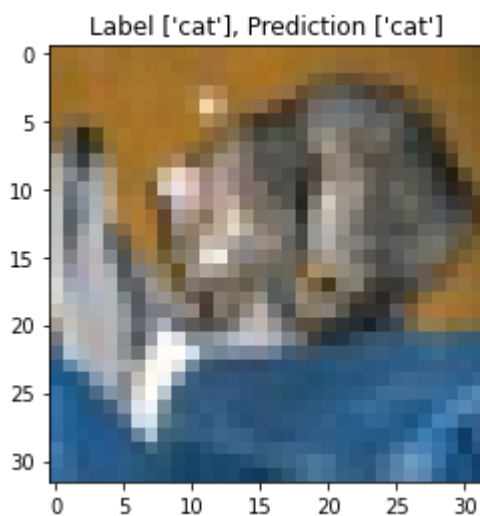
In terms of observations, it seems that the Maxpool and Original method seem to have similar trends for the accuracy and loss rate. However, it appears that Maxpool method is slightly less accurate than the original method.

```
In [14]: #To-do: evaluate on test set, instead of training set
random_image = random.randint(0,len(test_dataset))
image = test_dataset.__getitem__(random_image)
model.eval()
images, labels = next(iter(test_loader))
images, labels = images.to(device), labels.to(device)
predictions = torch.argmax(model(images),1)
num_cols = 1
num_rows = 25# Len(labels)
label_map = [['airplane'], ['automobile'], ['bird'], ['cat'], ['deer'], ['dog'],

for idx in range(num_rows):
    img = images.cpu()[idx]

    plt.title(f"Label {label_map[labels[idx]]}, Prediction {label_map[predictions[idx]]}")
    imshow(img)

    plt.axis('off')
plt.show()
```



ii.)

```

In [15]: class CNNModel(nn.Module):
def __init__(self):
    super(CNNModel, self).__init__()
    # TODO: Create CNNModel using 2D convolution. You should vary the number of
    # In this function, you should define each of the individual components of
    # Example:
    self.cnn1 = torch.nn.Conv2d(in_channels = 3, out_channels = 6, kernel_size
    self.relu = torch.nn.ReLU()
    #Input = 6 x 28 x 28, Output = 16 x 26 x 26
    self.cnn2 = torch.nn.Conv2d(in_channels = 6, out_channels = 16, kernel_size

    self.maxpool1 = nn.MaxPool2d(2)

    self.cnn3 = torch.nn.Conv2d(in_channels = 16, out_channels = 24, kernel_si
    self.cnn4 = torch.nn.Conv2d(in_channels = 24, out_channels = 24, kernel_si
    self.cnn5 = torch.nn.Conv2d(in_channels = 24, out_channels = 24, kernel_si

    # TODO: Create Fully connected layers. You should calculate the dimension
    # Example:
    # self.fc1 = nn.Linear(16 * 110 * 110, 5)
    # Fully connected 1
    self.fc1 = torch.nn.Linear(24*7*7, 120)
    self.fc2 = torch.nn.Linear(120, 84)
    self.fc3 = torch.nn.Linear(84, 10)

def forward(self,x):

    # TODO: Perform forward pass in below section
    # In this function, you will apply the components defined earlier to the i
    # Example:
    out = self.cnn1(x)
    out = self.relu(out)
    out = self.cnn2(out)
    out = self.relu(out)
    #plt.imshow(out[0][0].cpu().detach().numpy())
    #plt.show()
    #plt.close('all')
    out = self.maxpool1(out)
    out = self.cnn3(out)
    out = self.relu(out)
    out = self.cnn4(out)
    out = self.relu(out)
    out = self.cnn5(out)
    out = self.relu(out)
    #out = self.relu1(out)

    # to visualize feature map in part a, part b.i), use the following three l

    out = out.view(out.size(0), -1)
    out = self.fc1(out)
    out = self.relu(out)
    out = self.fc2(out)

```

```
        out = self.relu(out)
        out = self.fc3(out)

    return out

# Create CNN
device = "cuda" if torch.cuda.is_available() else "cpu"
model = CNNModel()
model.to(device)

# TODO: define Cross Entropy Loss
error = torch.nn.CrossEntropyLoss()

# TODO: create Adam Optimizer and define your hyperparameters
learning_rate = 1e-3
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
num_epochs = 20

from torchsummary import summary
print(summary(model, input_size=(3, 32, 32)))

count = 0
loss_list_bii = []
iteration_list_bii = []
accuracy_list_bii = []
for epoch in tqdm(range(num_epochs)):
    model.train()
    for i, (images, labels) in enumerate(train_loader):
        images, labels = images.to(device), labels.to(device)

        # Clear gradients
        optimizer.zero_grad()

        # TODO: Forward propagation
        outputs = model(images)

        # TODO: Calculate softmax and cross entropy Loss
        loss = error(outputs, labels)

        # Backpropagate your Loss
        T = torch.sum(loss)
        T.backward()

        # Update CNN model
        optimizer.step()
        count = count + 1

    if count % 50 == 0:
        model.eval()
        # Calculate Accuracy
        correct = 0
        total = 0
        # Iterate through test dataset
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
```

```

# Forward propagation
outputs = model(images)

# Get predictions from the maximum value
predicted = torch.argmax(outputs,1)

# Total number of labels
total = total + len(labels)

correct = correct + (predicted == labels).sum()

accuracy = 100 * correct / float(total)

# store loss and iteration
loss_list_bii.append(loss.item())
iteration_list_bii.append(count)
accuracy_list_bii.append(accuracy.item())
if count % 500 == 0:
    # Print Loss
    print('Iteration: {} Loss: {} Accuracy: {} %'.format(count, loss

```

```

=====
Layer (type:depth-idx)          Param #
=====
|---Conv2d: 1-1                  456
|---ReLU: 1-2                   --
|---Conv2d: 1-3                 880
|---MaxPool2d: 1-4              --
|---Conv2d: 1-5                 3,480
|---Conv2d: 1-6                 5,208
|---Conv2d: 1-7                 5,208
|---Linear: 1-8                 141,240
|---Linear: 1-9                 10,164
|---Linear: 1-10                850
=====
Total params: 167,486
Trainable params: 167,486
Non-trainable params: 0
=====
Layer (type:depth-idx)          Param #
=====
|---Conv2d: 1-1                  456
|---ReLU: 1-2                   --
|---Conv2d: 1-3                 880
|---MaxPool2d: 1-4              --
|---Conv2d: 1-5                 3,480
|---Conv2d: 1-6                 5,208
|---Conv2d: 1-7                 5,208
|---Linear: 1-8                 141,240
|---Linear: 1-9                 10,164
|---Linear: 1-10                850
=====
Total params: 167,486
Trainable params: 167,486
Non-trainable params: 0

```



```
5%|█          | 1/20 [00:31<10:01, 31.66s/it]
Iteration: 500  Loss: 1.5793367624282837  Accuracy: 40.55999755859375 %
10%|██         | 2/20 [01:07<10:13, 34.07s/it]
Iteration: 1000 Loss: 1.415574073791504  Accuracy: 47.34000015258789 %
15%|███        | 3/20 [01:42<09:44, 34.39s/it]
Iteration: 1500 Loss: 1.2583825588226318  Accuracy: 53.87999725341797 %
25%|████       | 5/20 [02:51<08:39, 34.63s/it]
Iteration: 2000 Loss: 1.2385494709014893  Accuracy: 55.21999740600586 %
30%|█████      | 6/20 [03:20<07:36, 32.60s/it]
Iteration: 2500 Loss: 1.1462451219558716  Accuracy: 58.21999740600586 %
35%|██████     | 7/20 [03:50<06:51, 31.62s/it]
Iteration: 3000 Loss: 1.3434181213378906  Accuracy: 58.55999755859375 %
40%|███████    | 8/20 [04:23<06:27, 32.28s/it]
Iteration: 3500 Loss: 0.960294783115387  Accuracy: 60.349998474121094 %
50%|████████   | 10/20 [05:30<05:26, 32.68s/it]
Iteration: 4000 Loss: 0.9490000009536743  Accuracy: 58.87999725341797 %
55%|█████████  | 11/20 [06:04<04:57, 33.04s/it]
Iteration: 4500 Loss: 0.9997236132621765  Accuracy: 60.18000030517578 %
60%|██████████ | 12/20 [06:36<04:22, 32.79s/it]
Iteration: 5000 Loss: 0.9303755164146423  Accuracy: 62.22999954223633 %
70%|███████████| 14/20 [07:45<03:21, 33.65s/it]
Iteration: 5500 Loss: 0.8039083480834961  Accuracy: 63.18000030517578 %
75%|███████████| 15/20 [08:20<02:50, 34.13s/it]
Iteration: 6000 Loss: 0.7739051580429077  Accuracy: 63.029998779296875 %
80%|███████████| 16/20 [08:56<02:18, 34.59s/it]
Iteration: 6500 Loss: 1.0581154823303223  Accuracy: 62.97999954223633 %
85%|███████████| 17/20 [09:29<01:42, 34.22s/it]
Iteration: 7000 Loss: 0.8343865871429443  Accuracy: 63.349998474121094 %
95%|███████████| 19/20 [10:39<00:34, 34.81s/it]
Iteration: 7500 Loss: 0.9185515642166138  Accuracy: 62.66999816894531 %
100%|███████████| 20/20 [11:15<00:00, 33.77s/it]
```



```
In [16]: # visualization loss
plt.plot(iteration_list_bi,loss_list_bi, color = "blue", label = 'Maxpool')
plt.plot(iteration_list_bii,loss_list_bii, color = "red", label = 'Five Convolution')
plt.xlabel("Number of iteration")
plt.ylabel("Loss")
plt.title("CNN: Loss vs Number of iteration")
plt.legend()
plt.show()

# visualization accuracy
plt.plot(iteration_list_bii,accuracy_list_bii,color = "red", label = 'Five Con')
plt.plot(iteration_list_bi,accuracy_list_bi,color = "blue", label = 'Maxpool')
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("CNN: Accuracy vs Number of iteration")
plt.legend()
plt.show()

print("\nMemory Required for Model: 1.00 MB")
```



Memory Required for Model: 1.00 MB

```
In [17]: print("As the method in part 3 uses more training hyper parameters than 2, we
```

As the method in part 3 uses more training hyper parameters than 2, we can see that it tends to be more accurate and takes up much less memory as well relative to the older method. However, the drawback to this is that we have a decrease in loss compared to older methods due to the more channels used to counter hyperparameters.