

## Note for question3

- Please follow the template to complete q3
- You may create new cells to report your results and observations

```
In [1]: # Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```


## P1. Load data and plot

### TODO

- load q3\_data.csv
- plot the points of different labels with different color

```
In [2]: # Load dataset
data = pd.read_csv("q3_data.csv", header = None)
y = np.array(data[2].astype(int))
# Plot points
plot = plt.scatter(data[0], data[1], c = y)
plt.legend(handles = plot.legend_elements()[0], labels = [0,1], title = "class")
plt.ylabel('x2')
plt.xlabel('x1')
```

Out[2]: Text(0.5, 0, 'x1')



## P2. Feature mapping

### TODO

- implement function **map\_feature()** to transform data from original space to the 28D space specified in the write-up

```
In [3]: # Transform points to 28D space
def map_feature(data):
    x1 = np.array(data[0])
```

```

x2 = np.array(data[1])
return np.array([x1*0 +1 , x1, x2, x1**2 , x1*x2, x2**2 , x1**3,
                  x1**2 * x2 , x1*x2**2 , x2**3, x1**4, x1**3 * x2,
                  x1**2 * x2 **2,x1 * x2 **3,x2 **4, x1**5, x1**4 * x2,
                  x1**3 * x2 **2,x1**2 * x2 **3,x1 * x2 **4, x2 **5,
                  x1**6, x1**5 * x2**1,x1**4 * x2**2,x1**3 * x2**3,
                  x1**2 * x2**4, x1**1 * x2**5, x2**6])

feature_mapping = map_feature([data[0], data[1]])

```

## P3. Regularized Logistic Regression

### TODO

- implement function **logistic\_regression\_regularized()** as required in the write-up
- draw the decision boundary

### Hints

- recycling code from HW2 is allowed
- you may use functions defined this section for part 4 below
- although optional for the report, plotting the convergence curve will be helpful

```

In [4]: # Define your functions here
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def calculate_gradients(X, y, y_pred):
    return np.dot(X, y_pred - y.reshape(len(y_pred), 1)) / X.shape[0]

def update_weights(prev_weights, current_grads, lr):
    return prev_weights - lr * current_grads

def main(X, Y, w, l, learning_rate = 0.00005, num_steps = 50000):
    n = 28
    for i in range(num_steps):
        y_pred = sigmoid(np.dot(X.transpose(), w))
        current_grad = calculate_gradients(X, Y, y_pred)
        w = update_weights(w, current_grad + (1 * w) / n, learning_rate)
    return w

def predict(weights, X, Y, l):
    new_Y = sigmoid(X.transpose() @ main(X, Y, weights, l))
    return new_Y, main(X, Y, weights, l)

def logistic_regression_regularized(df , l):
    X = map_feature([df[0],df[1]])
    Y = np.array(df[2])
    init_weights = np.zeros((28,1))
    new_Y, w = predict(init_weights, X, Y,l)
    return new_Y, w

_, weights = logistic_regression_regularized(data,l)
print("Lambda: 1 Weights", weights)

# Plot decision boundary
def plot(Y):
    label = np.array(data[2].astype(int))

```

```

plt.scatter(data[0],data[1],c=label)
plt.ylabel('x2')
plt.xlabel('x1')
plt.legend(handles = plt.scatter(data[0],data[1],c=label).legend_elements()[0], label='x2')

x_min = data[0].min()
x_max = data[0].max()
y_min = data[1].min()
y_max = data[1].max()
x_grid, y_grid = np.meshgrid(np.linspace(x_min, x_max,500), np.linspace(y_min, y_max,500))
xx = x_grid.ravel()
yy = y_grid.ravel()

p = np.dot(Y[1].transpose(),map_feature(np.vstack((xx,yy))))
p = p.reshape(x_grid.shape)
plt.contour(x_grid, y_grid, p, levels = [.5])
plt.show()

log_reg = logistic_regression_regularized(data,1)

plot(log_reg)


```

Lambda: 1 Weights [[ 1.05729526]

```

[ 0.16204655]
[ 0.26984668]
[ 0.15784554]
[-0.14896965]
[ 0.25385408]
[ 0.27565389]
[ 0.18895441]
[ 0.10681437]
[ 0.10799886]
[-0.09834135]
[ 0.0589715 ]
[-0.00956618]
[-0.04051705]
[-0.36657207]
[ 0.41416441]
[ 0.16261339]
[ 0.12956978]
[ 0.04589712]
[ 0.17802265]
[ 0.13711344]
[-0.21389026]
[-0.02530485]
[-0.2587769 ]
[ 0.05417648]
[-0.08400874]
[-0.06082876]
[-0.01638283]]

```




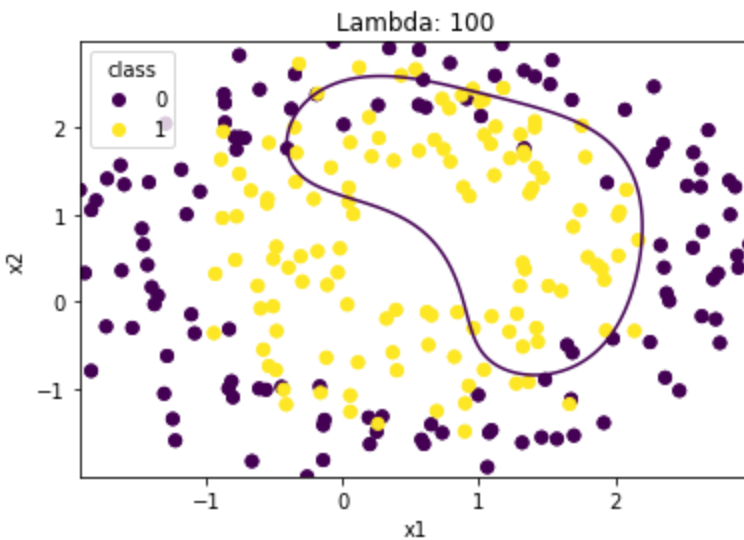
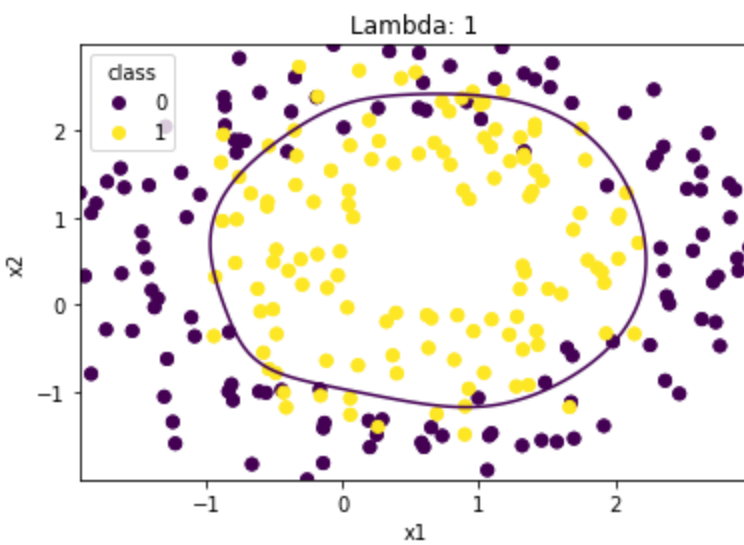
## P4. Tune the strength of regularization

### TODO

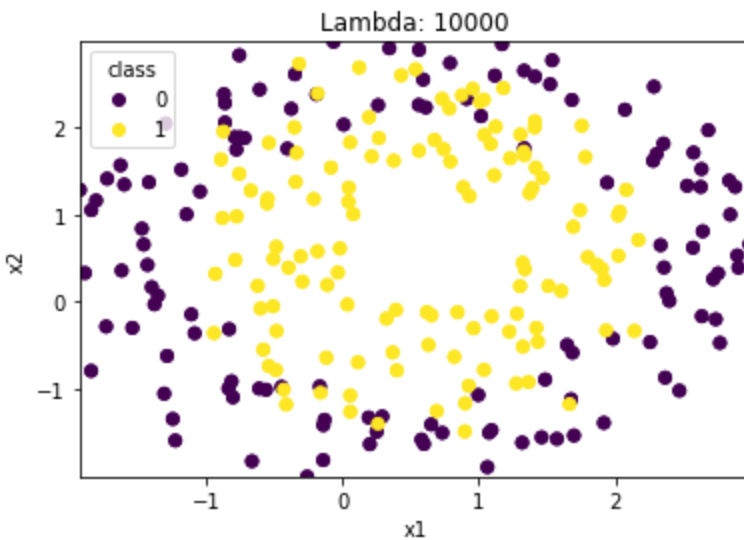
- tweak the hyper-parameter  $\lambda$  to be  $[0, 1, 100, 10000]$
- draw the decision boundaries

```
In [5]: # lambda = 0
lam_0 = logistic_regression_regularized(data, 0)
plt.title("Lambda: 0")
plot(lam_0)
# lambda = 1
lam_1 = logistic_regression_regularized(data, 1)
plt.title("Lambda: 1")
plot(lam_1)
# lambda = 100
lam_100 = logistic_regression_regularized(data, 100)
plt.title("Lambda: 100")
plot(lam_100)
# lambda = 10000
lam_10000 = logistic_regression_regularized(data, 10000)
plt.title("Lambda: 10000")
plot(lam_10000)
```





C:\Users\gandi\AppData\Local\Temp\ipykernel\_45664\1163830579.py:53: UserWarning: No contour levels were found within the data range.  
 plt.contour(x\_grid, y\_grid, p, levels = [.5])



Answer for part (d) here: As the value of lambda increases, we notice that the decision boundary seems to tighten up and become less accurate as well. There are many more false classifications as 1 instead due to the weights not being changed significantly at the higher lambda values. The lower lambda values seem to be a better fit as the weights adjust properly at each cycle whereas the model is underfit at the high lambdas.